

A Self-correcting Graph Connected Component Algorithm

Piyush Sao, Oded Green, Chirag Jain, Richard Vuduc

The 6th Workshop on Fault Tolerance for HPC at eXtreme Scale
(FTXS) 2016

<http://hpcgarage.org/ftxs16/>



Summary of Contributions

Self-correcting Algorithms

We introduce a new fault tolerant algorithm design principle that we call *self-correction*. A self-correcting algorithm remains in a valid state, despite the faulty execution of an iteration, under the assumption that its previous state was a valid one.

Self-Correcting Connected Components Algorithm

- We apply the ideas of self-correction to Label-propagation algorithm for graph connected component problem.
- Assumes availability of selective reliability mode.
- Requires $\mathcal{O}(V)$ additional storage and computations per iteration compared to $\mathcal{O}(|V| + |E|)$ cost for the baseline algorithm.
- 10-35% increases in execution time for one error for 64 memory operations.

Summary of Contributions


Self-correcting Algorithms

We introduce a new fault tolerant algorithm design principle that we call *self-correction*. A self-correcting algorithm remains in a valid state, despite the faulty execution of an iteration, under the assumption that its previous state was a valid one.

Self-Correcting Connected Components Algorithm

- We apply the ideas of self-correction to Label-propagation algorithm for graph connected component problem.
- Assumes availability of selective reliability mode.
- Requires $\mathcal{O}(V)$ additional storage and computations per iteration compared to $\mathcal{O}(|V| + |E|)$ cost for the baseline algorithm.
- 10-35% increases in execution time for one error for 64 memory operations.


Iterative Algorithms



Iterative Algorithms

- A typical iterative algorithm has following components:
 - 1 The input problem;
 - 2 Intermediate variable;
 - 3 Update relation;
 - 4 Convergence checking; and
 - 5 Solution.

Iterative Algorithm as State Machine




Iterative Algorithms as State Machines

- An iterative algorithm can be viewed as state machine.
- State of the algorithm: subset of intermediate variables required for continued execution of the algorithm.
- Starts with an initial state S_0
- Uses update relation to transition from one state to another

$$S_{k+1} \leftarrow S_k$$

Iterative Algorithm as State Machine




Iterative Algorithms as State Machines

- An iterative algorithm can be viewed as state machine.
- State of the algorithm: subset of intermediate variables required for continued execution of the algorithm.
- Starts with an initial state S_0
- Uses update relation to transition from one state to another

$$S_{k+1} \leftarrow S_k$$

Iterative Algorithm as State Machine




Iterative Algorithms as State Machines

- An iterative algorithm can be viewed as state machine.
- State of the algorithm: subset of intermediate variables required for continued execution of the algorithm.
- Starts with an initial state S_0
- Uses update relation to transition from one state to another

$$S_{k+1} \leftarrow S_k$$

Iterative Algorithm as State Machine




Iterative Algorithms as State Machines

- An iterative algorithm can be viewed as state machine.
- State of the algorithm: subset of intermediate variables required for continued execution of the algorithm.
- Starts with an initial state S_0
- Uses update relation to transition from one state to another

$$S_{k+1} \leftarrow S_k$$


Single Fault in Iterative Algorithm



Valid and Invalid States

- Valid state: under fault-free execution of the algorithm from that state, the algorithm will converge to the correct solution; otherwise invalid.
- In fault free execution, the algorithm always remains in a valid state.
- Any hardware fault can cause the algorithm to reach an invalid state.
- In general determining whether a given state is valid or not, is non-trivial.


Single Fault in Iterative Algorithm



Valid and Invalid States

- Valid state: under fault-free execution of the algorithm from that state, the algorithm will converge to the correct solution; otherwise invalid.
- In fault free execution, the algorithm always remains in a valid state.
- Any hardware fault can cause the algorithm to reach an invalid state.
- In general determining whether a given state is valid or not, is non-trivial.


Single Fault in Iterative Algorithm



Valid and Invalid States

- Valid state: under fault-free execution of the algorithm from that state, the algorithm will converge to the correct solution; otherwise invalid.
- In fault free execution, the algorithm always remains in a valid state.
- Any hardware fault can cause the algorithm to reach an invalid state.
- In general determining whether a given state is valid or not, is non-trivial.

Self-stabilizing Algorithms



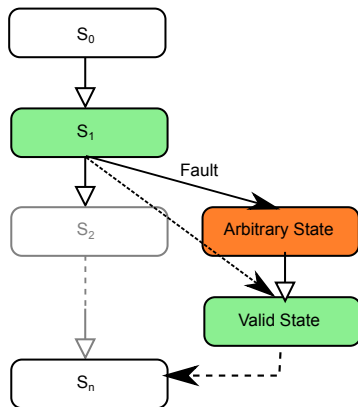
Self-stabilizing Algorithms

- Starting from any arbitrary state, valid or invalid, a self-stabilizing algorithm will reach a valid in finite number of steps.
- Natural fault-tolerance mechanism.
- Examples: Stationary iterations, Newton Iteration.
- Self-stabilization is a *strong* property.

Scala'13

- Self-stabilizing Steepest Descent and Conjugate Gradient.
- Periodic state correction.
- May not be generalized to all iterative algorithms.

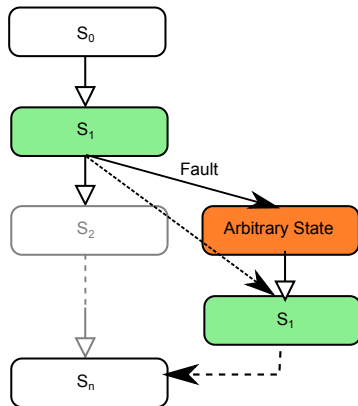
Self-Correcting Algorithms



Self-correcting Algorithms

- A self-correcting algorithm is an iterative algorithm that, starting in some valid state, remains in a valid state or comes to a valid state in finite number of steps even if a fault occurs.
- Requires that algorithm starts from a valid state.
- Uses information from previously known valid state.
- Example: Checkpoint-restart, FT-GMRES.

Checkpoint-restart as a Self-correcting algorithm



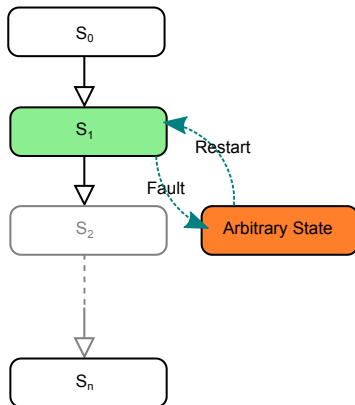
Checkpoint-restart based fault tolerance

- Bring to valid state by restoring a check-pointed valid state.
- At high fault rate, algorithm will not make any progress.

Broader idea of self-correction is to use S_1 to construct an state

$$\tilde{S}_2 \approx S_2$$

Checkpoint-restart as a Self-correcting algorithm



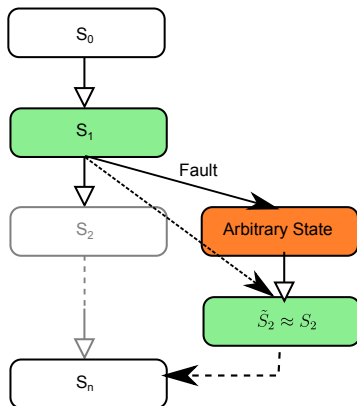
Checkpoint-restart based fault tolerance

- Bring to valid state by restoring a checkpointed valid state.
- **At high fault rate, algorithm will not make any progress.**

Broader idea of self-correction is to use S_1 to construct an state

$$\tilde{S}_2 \approx S_2$$

Checkpoint-restart as a Self-correcting algorithm




Checkpoint-restart based fault tolerance

- Bring to valid state by restoring a check-pointed valid state.
- At high fault rate, algorithm will not make any progress.

Broader idea of self-correction is to use S_1 to construct an state

$$\tilde{S}_2 \approx S_2$$

Label Propagation Algorithm for Graph Connected Component Algorithm




Graph Connected-component Problem

We seek to find number of connected-components in the graph and which connected component each vertex belongs to.

- Used for community detection, centrality analytics and streaming graph analytics.
- Label propagation is highly suited for parallel computing.

Label Propagation Algorithm for Graph Connected Component Algorithm




Graph Connected-component Problem

We seek to find number of connected-components in the graph and which connected component each vertex belongs to.

- Used for community detection, centrality analytics and streaming graph analytics.
- Label propagation is highly suited for parallel computing.

Label Propagation Algorithm




- Propagates the minimum label in the connected component.
- We maintain a *label* array CC for each vertex.
- CC is initialized to vertex id for each vertex.
- In each iteration, each vertex calculates minimum label among all near-neighbours
 $\mathcal{N}(v) = v \cup \text{adj}(v)$

$$CC^i[v] = \min_{u \in \mathcal{N}(v)} CC^{i-1}[u].$$

- Iteration terminates when no change occur in an iteration.

Label Propagation Algorithm




- Propagates the minimum label in the connected component.
- We maintain a *label* array CC for each vertex.
- CC is initialized to vertex id for each vertex.
- In each iteration, each vertex calculates minimum label among all near-neighbours
 $\mathcal{N}(v) = v \cup adj(v)$

$$CC^i[v] = \min_{u \in \mathcal{N}(v)} CC^{i-1}[u].$$


- Iteration terminates when no change occur in an iteration.

Label Propagation Algorithm- Example




- Minimum label in the connected component (shown in green and orange) propagates to all the vertex in the connected components.

Label Propagation Algorithm- Example




- Minimum label in the connected component (shown in green and orange) propagates to all the vertex in the connected components.

Label Propagation Algorithm- Example




- Minimum label in the connected component (shown in green and orange) propagates to all the vertex in the connected components.

Label Propagation Algorithm- Example




- Minimum label in the connected component (shown in green and orange) propagates to all the vertex in the connected components.

State of Label Propagation Algorithm




- CC vector defines the state of the LP algorithm.

Single Fault In Label Propagation Algorithm- Fault Correction




- CC value can be corrupted due to hardware faults.
- Depending on error caused by the fault, some error can be corrected by the algorithm.
- Example: corrupted $CC[v]$ value is arbitrarily high.
- Such faults do not cause the algorithm to transition to an invalid state, but they can still cause delay in convergence.

Single Fault In Label Propagation Algorithm- Fault Correction




- CC value can be corrupted due to hardware faults.
- Depending on error caused by the fault, some error can be corrected by the algorithm.
- Example: corrupted $CC[v]$ value is arbitrarily high.
- Such faults do not cause the algorithm to transition to an invalid state, but they can still cause delay in convergence.

Single Fault In Label Propagation Algorithm- Fault Propagation




- If the fault causes a corruption such $CC[v]$ is changed to a values lower than the $CC^\infty[v]$, the error will propagate to all the other vertex in the connected component.
- Such faults do not cause the algorithm to transition to an invalid state.
- It is not computationally easy to determine whether a given snapshot of CC is a valid state.

Single Fault In Label Propagation Algorithm- Fault Propagation



- If the fault causes a corruption such $CC[v]$ is changed to a values lower than the $CC^\infty[v]$, the error will propagate to all the other vertex in the connected component.
- Such faults do not cause the algorithm to transition to an invalid state.
- It is not computationally easy to determine whether a given snapshot of CC is a valid state.

Single Fault In Label Propagation Algorithm- Fault Propagation



- If the fault causes a corruption such $CC[v]$ is changed to a values lower than the $CC^\infty[v]$, the error will propagate to all the other vertex in the connected component.
- Such faults do not cause the algorithm to transition to an invalid state.
- It is not computationally easy to determine whether a given snapshot of CC is a valid state.

Valid States

Classification of different corruption

Consider following three cases:

- 1 $CC[v] > v$: Easy to detect and automatically corrected in most cases.
- 2 $CC^\infty[v] \leq CC[v] \leq v$: ??
- 3 $CC[v] < CC^\infty[v]$: Will definitely cause algorithm to fail.

Theorem

A connected component array CC is a valid state—i.e., a fault-free execution of algorithm starting from CC will converge to the correct solution—if, for all vertices v ,

$$CC^\infty[v] \leq CC[v] \leq v.$$

The only problem is, $CC^\infty[v]$ —being the solution, is not known.

Valid States

Classification of different corruption

Consider following three cases:

- 1 $CC[v] > v$: Easy to detect and automatically corrected in most cases.
- 2 $CC^\infty[v] \leq CC[v] \leq v$: ??
- 3 $CC[v] < CC^\infty[v]$: Will definitely cause algorithm to fail.

Theorem

A connected component array CC is a valid state—i.e., a fault-free execution of algorithm starting from CC will converge to the correct solution—if, for all vertices v ,

$$CC^\infty[v] \leq CC[v] \leq v.$$

The only problem is, $CC^\infty[v]$ —being the solution, is not known.

Valid States

Classification of different corruption

Consider following three cases:

- 1 $CC[v] > v$: Easy to detect and automatically corrected in most cases.
- 2 $CC^\infty[v] \leq CC[v] \leq v$: ??
- 3 $CC[v] < CC^\infty[v]$: Will definitely cause algorithm to fail.

Theorem

A connected component array CC is a valid state—i.e., a fault-free execution of algorithm starting from CC will converge to the correct solution—if, for all vertices v ,

$$CC^\infty[v] \leq CC[v] \leq v.$$

The only problem is, $CC^\infty[v]$ —being the solution, is not known.

Valid States

Classification of different corruption

Consider following three cases:

- 1 $CC[v] > v$: Easy to detect and automatically corrected in most cases.
- 2 $CC^\infty[v] \leq CC[v] \leq v$: ??
- 3 $CC[v] < CC^\infty[v]$: Will definitely cause algorithm to fail.

Theorem

A connected component array CC is a valid state—i.e., a fault-free execution of algorithm starting from CC will converge to the correct solution—if, for all vertices v ,

$$CC^\infty[v] \leq CC[v] \leq v.$$

The only problem is, $CC^\infty[v]$ —being the solution, is not known.

Valid States

Classification of different corruption

Consider following three cases:

- 1 $CC[v] > v$: Easy to detect and automatically corrected in most cases.
- 2 $CC^\infty[v] \leq CC[v] \leq v$: ??
- 3 $CC[v] < CC^\infty[v]$: Will definitely cause algorithm to fail.

Theorem

A connected component array CC is a valid state—i.e., a fault-free execution of algorithm starting from CC will converge to the correct solution—if, for all vertices v ,

$$CC^\infty[v] \leq CC[v] \leq v.$$

The only problem is, $CC^\infty[v]$ —being the solution, is not known.

Self-correcting Label Propagation Algorithm- 1

We apply principle of self-correction to resolve the apparent difficulty in verifying state validity.

- We assume the previous state CC^{i-1} is a valid one.
- Checking

$$CC^i[v] = \min_{u \in \mathcal{N}(v)} CC^{i-1}[u],$$

where $\mathcal{N}(v) \equiv v \cup \text{adj}(v)$ is the immediate *neighborhood* of v , **will require re-computing entire iteration.**

- We show that $CC^i[v]$ is still a valid value even if we can relax the minimization criterion to

$$CC^i[v] \in \{CC^{i-1}[u] \mid u \in \mathcal{N}(v)\},$$

Self-correcting Label Propagation Algorithm- 1

We apply principle of self-correction to resolve the apparent difficulty in verifying state validity.

- We assume the previous state CC^{i-1} is a valid one.
- Checking

$$CC^i[v] = \min_{u \in \mathcal{N}(v)} CC^{i-1}[u],$$

where $\mathcal{N}(v) \equiv v \cup \text{adj}(v)$ is the immediate *neighborhood* of v , **will require re-computing entire iteration.**

- We show that $CC^i[v]$ is still a valid value even if we can relax the minimization criterion to

$$CC^i[v] \in \{CC^{i-1}[u] \mid u \in \mathcal{N}(v)\},$$

Self-correcting Label Propagation Algorithm- 1

We apply principle of self-correction to resolve the apparent difficulty in verifying state validity.

- We assume the previous state CC^{i-1} is a valid one.
- Checking

$$CC^i[v] = \min_{u \in \mathcal{N}(v)} CC^{i-1}[u],$$

where $\mathcal{N}(v) \equiv v \cup \text{adj}(v)$ is the immediate *neighborhood* of v , **will require re-computing entire iteration.**

- We show that $CC^i[v]$ is still a valid value even if we can relax the minimization criterion to

$$CC^i[v] \in \{CC^{i-1}[u] \mid u \in \mathcal{N}(v)\},$$

Self-correcting Label Propagation Algorithm- 1

We apply principle of self-correction to resolve the apparent difficulty in verifying state validity.

- We assume the previous state CC^{i-1} is a valid one.
- Checking

$$CC^i[v] = \min_{u \in \mathcal{N}(v)} CC^{i-1}[u],$$

where $\mathcal{N}(v) \equiv v \cup \text{adj}(v)$ is the immediate *neighborhood* of v , **will require re-computing entire iteration.**

- We show that $CC^i[v]$ is still a valid value even if we can relax the minimization criterion to

$$CC^i[v] \in \{CC^{i-1}[u] \mid u \in \mathcal{N}(v)\},$$

Self-correcting Label Propagation Algorithm- 2

Theorem

Given a valid state for the previous iteration, CC^{i-1} , the current connected component array CC^i is a valid state if for all vertices v , CC^i satisfies these conditions:

- 1 $CC^i[v] \leq v$; and
- 2 there exists a vertex u such that $CC^i[v] = CC^{i-1}[u]$ and $u \in \mathcal{N}(v)$.

Cost of Direct Verification

- Verifying $CC^i[v] \leq v$ requires $\mathcal{O}(V)$ operation.
- Verifying second condition requires traversing adjacency list for each vertex v , that will require $\mathcal{O}(V + E)$ operations, as costly as an LP iteration.

Self-correcting Label Propagation Algorithm- 2

Theorem


Given a valid state for the previous iteration, CC^{i-1} , the current connected component array CC^i is a valid state if for all vertices v , CC^i satisfies these conditions:

- ❶ $CC^i[v] \leq v$; and
- ❷ there exists a vertex u such that $CC^i[v] = CC^{i-1}[u]$ and $u \in \mathcal{N}(v)$.

Cost of Direct Verification

- Verifying $CC^i[v] \leq v$ requires $\mathcal{O}(V)$ operation.
- Verifying second condition requires traversing adjacency list for each vertex v , that will require $\mathcal{O}(V + E)$ operations, **as costly as an LP iteration.**

Validity Checking: Auxiliary Data Structure- 1




Parent Array

- Parent array P : We may store information of the vertex u that caused the last change in $CC[v]$.
- If $u = P[v]$ then $CC^i[v] = CC^{i-1}[P[v]]$, can be verified in $\mathcal{O}(V)$ operations for all vertex.
- Storing P requires a memory of $\mathcal{O}(V)$.

Corruption of P

- P also can be corrupt.
- P is valid if $P[v] \in \mathcal{N}(v)$ for all vertex v .
- Checking P is valid requires again $\mathcal{O}(V + E)$ operations.

Validity Checking: Auxiliary Data Structure- 1




Parent Array

- Parent array P : We may store information of the vertex u that caused the last change in $CC[v]$.
- If $u = P[v]$ then $CC^i[v] = CC^{i-1}[P[v]]$, can be verified in $\mathcal{O}(V)$ operations for all vertex.
- Storing P requires a memory of $\mathcal{O}(V)$.

Corruption of P

- P also can be corrupt.
- P is valid if $P[v] \in \mathcal{N}(v)$ for all vertex v .
- Checking P is valid requires again $\mathcal{O}(V + E)$ operations.

Validity Checking: Auxiliary Data Structure- 1




Parent Array

- Parent array P : We may store information of the vertex u that caused the last change in $CC[v]$.
- If $u = P[v]$ then $CC^i[v] = CC^{i-1}[P[v]]$, can be verified in $\mathcal{O}(V)$ operations for all vertex.
- Storing P requires a memory of $\mathcal{O}(V)$.

Corruption of P

- P also can be corrupt.
- P is valid if $P[v] \in \mathcal{N}(v)$ for all vertex v .
- Checking P is valid requires again $\mathcal{O}(V + E)$ operations.

Validity Checking: Auxiliary Data Structure- 2



Index Based Parent Array

- Instead of storing u , we store index of u in $adj(v)$.

$E \quad \leftarrow adj(v)$

$u \quad \leftarrow E[k]$

$P^*[v] \quad \leftarrow k$


- When $P[v] = v$, then $P^*[v] = -1$
- $CC^i[v] = CC^{i-1}[P[v]]$ reduces to

$$CC^i[v] = CC^{i-1}[adj(v)[P^*[v]]];$$

$\mathcal{O}(V)$ operations.

- Validity of P^* :
 $-1 \leq P^*[v] < |adj(v)|$; $\mathcal{O}(V)$ operations.

Validity Checking: Auxiliary Data Structure- 2



Index Based Parent Array

- Instead of storing u , we store index of u in $adj(v)$.

$$E \quad \leftarrow adj(v)$$

$$u \quad \leftarrow E[k]$$

$$P^*[v] \quad \leftarrow k$$


- When $P[v] = v$, then $P^*[v] = -1$
- $CC^i[v] = CC^{i-1}[P[v]]$ reduces to

$$CC^i[v] = CC^{i-1}[adj(v)[P^*[v]]];$$

$\mathcal{O}(V)$ operations.

- Validity of P^* :
 $-1 \leq P^*[v] < |adj(v)|$; $\mathcal{O}(V)$ operations.

Validity Checking: Auxiliary Data Structure- 2



Index Based Parent Array

- Instead of storing u , we store index of u in $adj(v)$.

$$E \quad \leftarrow adj(v)$$

$$u \quad \leftarrow E[k]$$

$$P^*[v] \quad \leftarrow k$$


- When $P[v] = v$, then $P^*[v] = -1$
- $CC^i[v] = CC^{i-1}[P[v]]$ reduces to

$$CC^i[v] = CC^{i-1}[adj(v)[P^*[v]]];$$

$\mathcal{O}(V)$ operations.

- Validity of P^* :
 $-1 \leq P^*[v] < |adj(v)|$; $\mathcal{O}(V)$ operations.

Validity Checking: Auxiliary Data Structure- 2



Index Based Parent Array

- Instead of storing u , we store index of u in $\text{adj}(v)$.

$E \leftarrow \text{adj}(v)$

$u \leftarrow E[k]$

$P^*[v] \leftarrow k$


- When $P[v] = v$, then $P^*[v] = -1$
- $CC^i[v] = CC^{i-1}[P[v]]$ reduces to

$$CC^i[v] = CC^{i-1}[\text{adj}(v)[P^*[v]]];$$

$\mathcal{O}(V)$ operations.

- Validity of P^* :
 $-1 \leq P^*[v] < |\text{adj}(v)|$; $\mathcal{O}(V)$ operations.

Validity Checking: Auxiliary Data Structure- 2



Index Based Parent Array

- Instead of storing u , we store index of u in $\text{adj}(v)$.

$$E \leftarrow \text{adj}(v)$$

$$u \leftarrow E[k]$$

$$P^*[v] \leftarrow k$$


- When $P[v] = v$, then $P^*[v] = -1$
- $CC^i[v] = CC^{i-1}[P[v]]$ reduces to

$$CC^i[v] = CC^{i-1}[\text{adj}(v)[P^*[v]]];$$

$\mathcal{O}(V)$ operations.

- Validity of P^* :
 $-1 \leq P^*[v] < |\text{adj}(v)|$; $\mathcal{O}(V)$ operations.

Validity Checking: Auxiliary Data Structure- 2



Index Based Parent Array

- Instead of storing u , we store index of u in $adj(v)$.

$$E \leftarrow adj(v)$$

$$u \leftarrow E[k]$$

$$P^*[v] \leftarrow k$$


- When $P[v] = v$, then $P^*[v] = -1$
- $CC^i[v] = CC^{i-1}[P[v]]$ reduces to

$$CC^i[v] = CC^{i-1}[adj(v)[P^*[v]]];$$

$\mathcal{O}(V)$ operations.

- Validity of P^* :
 $-1 \leq P^*[v] < |adj(v)|$; $\mathcal{O}(V)$ operations.

Validity Checking: Auxiliary Data Structure- 2



Index Based Parent Array

- Instead of storing u , we store index of u in $adj(v)$.

$$E \leftarrow adj(v)$$

$$u \leftarrow E[k]$$

$$P^*[v] \leftarrow k$$

- When $P[v] = v$, then $P^*[v] = -1$
- $CC^i[v] = CC^{i-1}[P[v]]$ reduces to

$$CC^i[v] = CC^{i-1}[adj(v)[P^*[v]]];$$

$\mathcal{O}(V)$ operations.

- Validity of P^* :
 $-1 \leq P^*[v] < |adj(v)|$; $\mathcal{O}(V)$ operations.

Fault Detection and Correction

Invalid State Detection

In summary, the set of conditions to check for each vertex are:

$$\begin{aligned} &CC^i[v] \leq v; \\ &-1 \leq P^*[v] < |adj(v)|; \text{ and} \\ &CC^i[v] = \begin{cases} v & \text{if } P^*[v] = -1 \\ CC^{i-1}[adj(v)[P^*[v]]] & \text{if } P^*[v] \neq -1 \end{cases} \end{aligned}$$

State Correction

For any vertex v , if state validity check fails then, we recompute $CC^i[v]$.

Fault Detection and Correction

Invalid State Detection

In summary, the set of conditions to check for each vertex are:

$$\begin{aligned} CC^i[v] &\leq v; \\ -1 &\leq P^*[v] < |adj(v)|; \text{ and} \\ CC^i[v] &= \begin{cases} v & \text{if } P^*[v] = -1 \\ CC^{i-1}[adj(v)[P^*[v]]] & \text{if } P^*[v] \neq -1 \end{cases} \end{aligned}$$

State Correction

For any vertex v , if state validity check fails then, we recompute $CC^i[v]$.

Overhead of Self-correcting Label-propagation Algorithm

Overhead	Asymptotic Complexity
Fault detection	$\mathcal{O}(V)$
Fault correction	$\mathcal{O}(f E / V)$
Auxiliary data structure	$\mathcal{O}(V)$

- Number of state corrections f , can be significantly less than faults occurred.
- Fault detection and correction needs to be done in a guaranteed reliable mode.

Experimental Setup

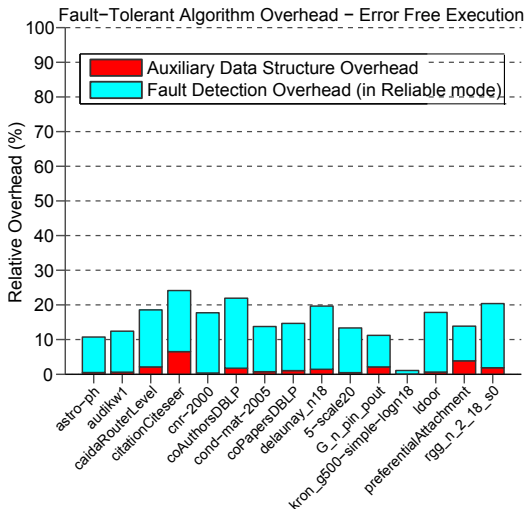
Machine Parameter

Prop	SNB16c
Micro-architecture	Sandy-Bridge
Sockets×Cores	2×8
Clock Rate	2.4GHz
DRAM capacity	128GB
DRAM Bandwidth	72GB/s
Compiler	Intel "C" compiler 15.0.0

Fault Injection


- Fault injection in reading graph data structure and CC array.
- Each fault injection read is independent
- Normalized by number of edges in the network
- Test Network: 14th DIMACS graph challenge

Fault Free Execution Overhead




On an average 1.3% overhead for maintaining additional data structure and 14% for fault detection.

Overhead of Fault Tolerant Algorithm in the Presence of Faults



Overhead of Fault Tolerant Algorithm in the Presence of Faults



Fault correction adds additional 9% overhead at 2^{-6} bit flips per every memory access.

Conclusion

Conclusion

- We introduced the ideas of self-correcting algorithm to build fault tolerant algorithms.
- We presented a self-correcting label propagation algorithm for graph connected component problem.
- Key steps involved:
 - Analyze valid and invalid state;
 - Use self-correction hypothesis to simplify invalid state detection;
 - Use previous valid states to recover from invalid state.
- Asymptotically lower overhead for fault detection and correction.
- 10-35% increases in execution time for one error for 64 memory operations.